

# *Database Systems:*

## *Module 12, Lecture 3 – The Relational Problem*

Instructor: Alan Paradise

# The Relational Problem

## LESSON OBJECTIVES

- To be able to describe the challenges that the explosion of "Big Data" brings to traditional **relational database** technology
- Describe various database technology solutions that organizations can implement to handle the explosion of "Big Data"

# The Relational Problem

## Situation:

- Big Data is exploding all around us, all over the world
- The explosion of Big Data is growing at an increasing rate
- Organizations struggle to cope with Big Data

## Problem:

40-year-old Relational Database technology cannot easily handle the volume, velocity and variety of Big Data.

# The Relational Problem

## The Relational Problem:

The Relational Database is 40+ year-old technology

- Demands STRUCTURE: Tables, Rows, Columns, Keys, Indexes
- Demands ACID Transaction Compliance
  - Keep my data consistent across transactions
  - Consistency VS Fast Performance and Throughput
- Prior RDBMS Software Choices
  - Investment in code
  - Investment in training staff

Traditional Database Server Architecture

- Memory, CPU, Storage
- The cost of scaling "UP" versus Scaling "OUT"

# The Relational Problem

## Questions:

My organization relies on operational software that sits on top of relational databases.

- How can I expand the capacity of my apps and their databases to handle our Big Data explosion?
- If I shift away from relational DBMS will I have to rewrite all my application code?
- If I shift away from relational DBMS will I have to retrain or replace all my software developers and support staff?
- Can I change my current systems (without replacing my databases) to expand capacity and run our queries faster?

# The Relational Problem

One approach:

- Keep my existing Relational DBMS systems and staff
- SCALE my systems to handle more data and run faster

I can SCALE my database servers

- UP: Adding CPU, Memory, Storage ("vertical")
- OUT: Add server nodes to a cluster ("horizontal")

# The Relational Problem

Scaling Up may be very costly

- Requires scalable, expandable database servers (\$\$\$)
- Requires purchasing more memory (\$\$\$)
- Requires purchasing more CPU capacity (\$\$\$)
- Requires purchasing more disk storage (\$\$)

# The Relational Problem

Scaling Out is a better solution

- Purchase cheap, commodity server hardware
- Connect multiple database servers together to form a "cluster"
  - Multiple database servers running a single database
  - Working together as a single server



# The Relational Problem

Google pioneered the concepts of scaling out clustered database servers

We can learn about clustering from Google's experience

# The Relational Problem

Be sure to watch the Google Container Data Center video. (2009, getting started with clustering)

<https://www.youtube.com/watch?v=zRwPSFpLX8I>

Challenges:

- Heat
- Maintenance
- UPS – Uninterruptible Power Supply

# The Relational Problem

Be sure to also watch the Google Data Center tours (9 years later)

<https://www.youtube.com/watch?v=zDAYZU4A3w0>

<https://www.youtube.com/watch?v=XZmGGAhHqa0>

<https://www.youtube.com/watch?v=avP5d16wEp0>

# The Relational Problem

In summary

- Google mastered horizontal scaling through clustering
- Cheap, commodity server hardware with directly attached disk storage
- Google figured out how to handle
  - server maintenance
  - heat
  - expandability

# The Relational Problem

Next Topic: Scaling Relational Systems